

Rechnerarchitektur 1

Inhaltsübersicht

- 1) Einführung, von Neumann Architektur
- 2) Pipeline-Prinzipien
 - a) Dependencies
 - b) Predication
- 3) Speicherhierarchie
 - a) Register
 - b) Caches
 - c) Hauptspeicher
 - d) Memory Management Unit (MMU)
- 4) Prozessorarchitekturen
 - a) Complex Instruction Set Computer (CISC)
Mikroprogrammierung
 - b) Reduced Instruction Set Computer (RISC)
- 5) Software Konzepte
 - a) Instruction Level Parallelism
 - b) Prozesse, Threads
 - c) Synchronisation
- 6) Einfache Verbindungsnetzwerke (Bus)
- 7) Moderne Prozessoren
 - a) Superscalar Processors
 - b) Very Long Instruction Word Machine (VLIW)
 - c) Multithreaded Architectures (MTA) -> shifted to RA2
 - d) Datenflußarchitekturen -> shifted to RA2
- 8) I/O Systeme
 - a) Prozessor-, System- & Peripheriebus
 - b) Ein-/Ausgabekomponenten
- 9) Netzwerk-Technologien (LAN/WAN)

Rechnerarchitektur 1

Literatur

Hennessy, John L.; Patterson, David A. (Standardwerk)

Computer architecture : A quantitative approach

2nd ed 1995. hardback £25.00

Morgan Kaufmann

1-55860-329-8

Giloi, Wolfgang K. (empfehlenswert, Übersicht)

Rechnerarchitektur

2.vollst. überarb. Aufl.1993 DM58.00

Springer-Verlag

3-540-56355-5

Hwang, Kai (advanced, viel parallele Arch.)

Advanced computer architecture : Parallelism, scalability and programmability

1993. 672 pp. paperback £24.99

McGraw-Hill ISE

0-07-113342-9

Patterson, David A.; Hennessy, John L. (leichter Einstieg)

Computer organization and design : The hardware/software interface

2nd ed 1994. hardback £60.50

Morgan Kaufmann

1-55860-281-X

Stone, Harold S. (optional, Speicherarchitekturen)

High performance computer architecture

3rd ed 1990. 736 pp. hardback £39.50

Addison-Wesley USA

0-201-52688-3

Tanenbaum, Andrew S. (optional, Standardwerk für Betriebssysteme))

Modern operating systems

1992. 1055 pp. paperback £26.95

Prentice Hall US

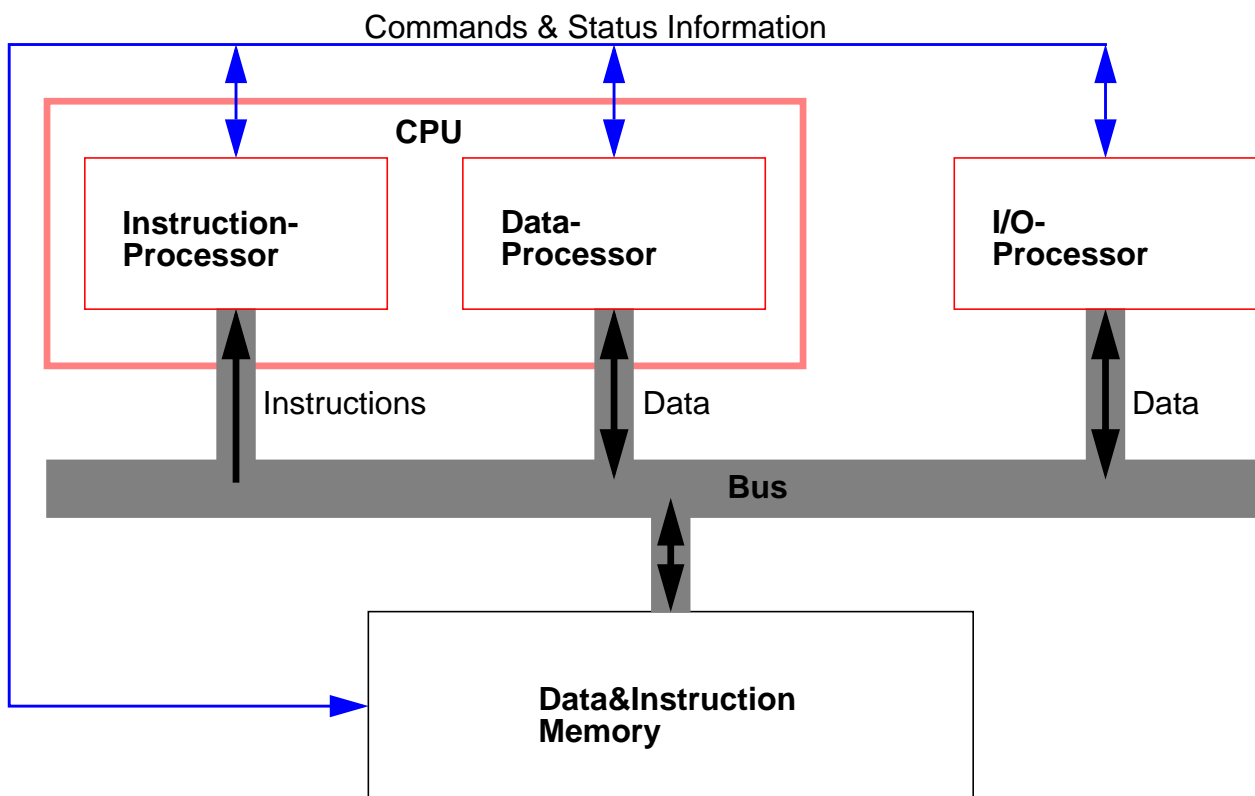
0-13-595752-4

Einleitung

von Neumann Architektur:

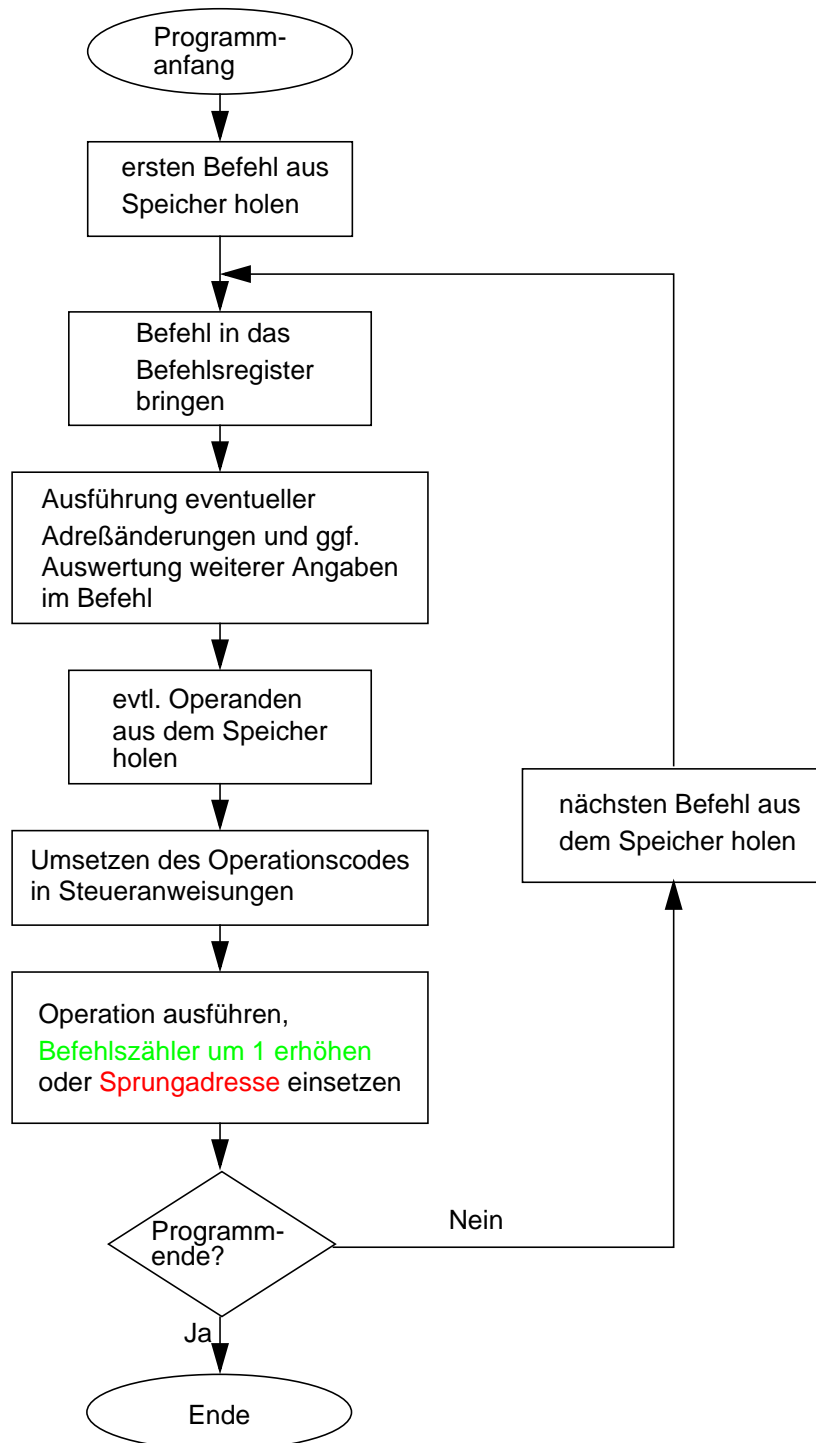
Beschreibt eine abstrakte Maschine des minimalen Hardwareaufwandes, bestehend aus:

- *einer zentralen Recheneinheit (CPU), die sich in Daten- und Befehlsprozessor aufteilt*
- *einem Speicher für Daten und Befehle*
- *einem Ein/Ausgabe-Prozessor zur Anbindung peripherer Geräte*
- *einer Verbindungseinrichtung (Bus) zwischen diesen Komponenten*



von Neumann Architektur

Verarbeitung von Instruktionen:



History of Computer Generations

History of computer generations

| Generation | Technology & Architecture | Software & Applikations | Systems |
|---------------------|---|---|--|
| First (1945-54) | Vacuum tubes and relay memories, CPU driven by PC and accumulator, fixed point arithmetic | Machine/assembly languages, single user, no subroutine linkage, programmed I/O using CPU | ENIAC, Princeton IAS, IBM 701 |
| Second (1955-64) | Discrete transistors and core memories, floating point arithmetic, I/O processors, multiplexed memory access | HLL used with compilers, subroutine libraries, batch processing monitor | IBM 7090, CDC 1604, Univac LARC |
| Third (1965-74) | Integrated circuits (S/MSI), microprogramming, pipelining, cache, look-ahead processors | Multiprogramming and timesharing OS, multiuser applications | IBM 360/70, CDC 6600, TI-ASC, PDP-8 |
| Fourth (1975-90) | LSI/VLSI and semiconductor memory, multiprocessors, vector supercomputers, multicomputers | Multiprocessor OS, languages, compilers and environment for parallel processing | VAX 9000, Cray X-MP, IBM 3090 |
| Fifth (1991-96) | ULSI/VHSIC processors, memory and switches, high density packaging | Massively parallel processing, grand challenge applications | Cray MPP, CM-5, Intel Paragon |
| Sixth (present) | scalable off-the-shelf architectures, workstation clusters, high speed interconnection networks | heterogeneous processing, fine grain message transfer | Gigabit Ethernet, Myrinet |

[Hwang, Advanced Computer Architecture]

Was ist Rechnerarchitektur ?

Rechnerarchitektur <- ' computer architecture '

The types of architecture are established not by architects but by society, according to the needs of the different institutions. Society sets the goals and assigns to the architect the job of finding the means of achieving them.

Entsprechend ihrem Zweck unterscheidet man in der Baukunst Typen von Architekturen

Arbeit des Computer-Architekten:

- finden eines Types von Architektur, die den vorgegebenen Zweck erfüllt.
- Architekt muß gewisse Forderungen erfüllen.
 - Leistung
 - Kosten
 - Ausfalltoleranz
 - Erweiterbarkeit

Materialien des Gebäudearchitekten : Holz, Stein, Beton, Glas ...

Materialien des Rechnerarchitekten : integrierte Halbleiterbausteine ...


Folgende Komponenten stellen die wesentlichen HW-Betriebsmittel des Computers dar:
Prozessoren, Speicher, Verbindungseinrichtungen

Vorschrift für das Zusammenfügen von Komponenten

- Operationsprinzip

Struktur der Anordnung von Komponenten

- Strukturanordnung



Bestandteile der Rechnerarchitektur

Eine Rechnerarchitektur ist bestimmt durch ein Operationsprinzip für die Hardware und die Struktur ihres Aufbaus aus den einzelnen Hardware-Betriebsmitteln. [Giloï 93]

Das Operationsprinzip

Das Operationsprinzip definiert das funktionelle Verhalten der Architektur durch Festlegung einer Informationsstruktur und einer Kontrollstruktur.

Die (Hardware) - Struktur

Die Struktur einer Rechnerarchitektur ist gegeben durch Art und Anzahl der Hardware-Betriebsmittel sowie die sie verbindenden Kommunikationseinrichtungen.

- Kontrollstruktur : Die Kontrollstruktur einer Rechnerarchitektur wird durch Spezifikation der Algorithmen für die Interpretation und Transformation der Informationskomponenten der Maschine bestimmt.
- Informationsstruktur : Die Informationsstruktur einer Rechnerarchitektur wird durch die Typen der Informationskomponenten in der Maschine bestimmt, der Repräsentation dieser Komponenten und der auf sie anwendbaren Operationen. Die Informationsstruktur läßt sich als Menge von 'abstrakten' Datentypen spezifizieren.
- Hardware-Betriebsmittel : Hardware-Betriebsmittel einer Rechnerarchitektur sind Prozessoren, Speicher, Verbindungseinrichtungen und Peripherie.
- Kommunikationseinrichtungen : Kommunikationseinrichtungen sind die Verbindungseinrichtungen, wie z.B.: Busse, Kanäle, VNs ; und die Protokolle, die die Kommunikationsregeln zwischen den Hardware-Betriebsmittel festlegen.

Unterscheidung der Parallelrechner nach dem Operationsprinzip

Unter dem Operationsprinzip versteht man das funktionelle Verhalten der Architektur, welches auf der zugrunde liegenden Informations- und Kontrollstruktur basiert.

Klassifikation nach Flynn (recht grob)

SISD single instruction - single data stream

SIMD single instruction - multiple data stream


MISD multiple instruction - single data stream

MIMD multiple instruction - multiple data stream

Unterscheidung der Parallelrechner nach dem Operationsprinzip

Beispiel : Automobilfertigung


Welche Möglichkeiten des Zusammenfügens (Assembly) gibt es ?



Worauf muß man achten ?

Abhängigkeiten !

Fließband



Produktion von verschiedenen Modellen : 3 Farben R(ot),
G(rün)
B(lau)

41

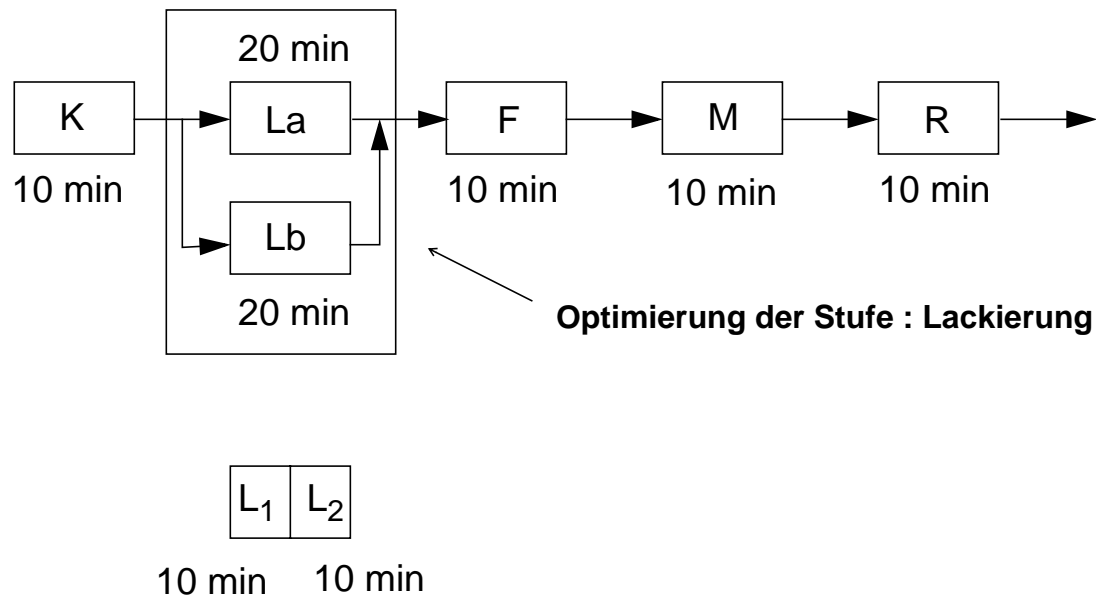
| Auftrag | |
|---------|---|
| F | R |
| M | N |
| K | K |

2 Motoren N(ormal),
I(njection)

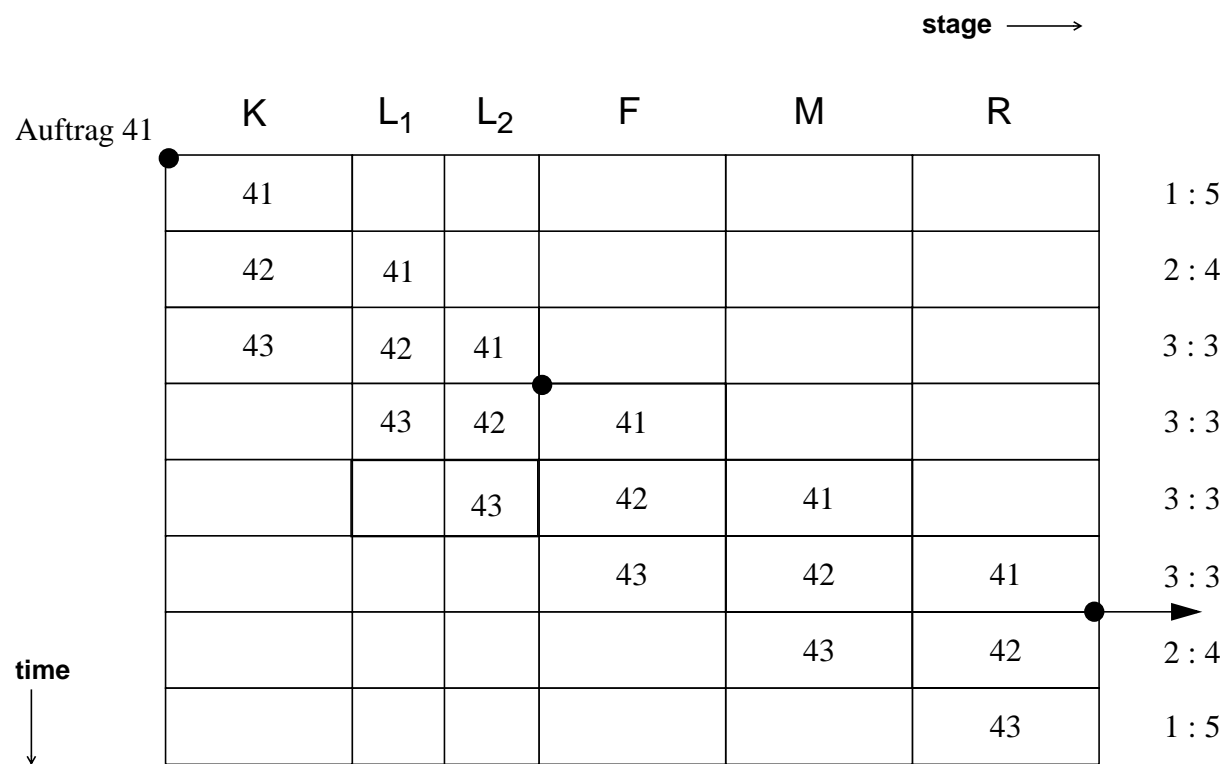
2 Karosserien L(imousine),
K(ombi)

Unterscheidung der Parallelrechner nach dem Operationsprinzip

Pipeline des Fertigungsvergangs



Stufen - Zeit - Diagramm der Pipeline



' von Neumann ' - Rechner (Burks, Goldstine, von Neumann)

Architektur des minimalen Hardwareaufwands !

Die hohen Kosten der Rechnerhardware in der Anfangszeit der Rechnerentwicklung erzwingen den möglichst kleinen Hardwareaufwand.

- Befehle holen
- Operand holen (2x) Indexoperation, Adressrechnung
- Befehle interpretieren
- Befehl ausführen arithmetische Op. (gr. Aufwand an Zeit + HW)

Ist die arithmetische Operation der wesentliche Aufwand, so können die anderen Teilschritte sequentiell ausgeführt werden, ohne die Verarbeitungszeit stark zu verlängern.

Zeitsequentielle Ausführung der Befehle auf minimalen Hardware-Ressourcen

Veränderung der Randbedingungen :

1. Reduzierung der Hardwarekosten
Hochintegration, Massenproduktion
2. Aufwand in den einzelnen Op. verschob sich.


Gesucht : Ideen, mit zusätzlicher Hardware bisher
 nicht mögliche Leistungssteigerungen
 zu erreichen.

Ideen = Operationsprinzipien

Technology

Introduction

First integrated microprocessor Intel 4004 1971 2700 transistors

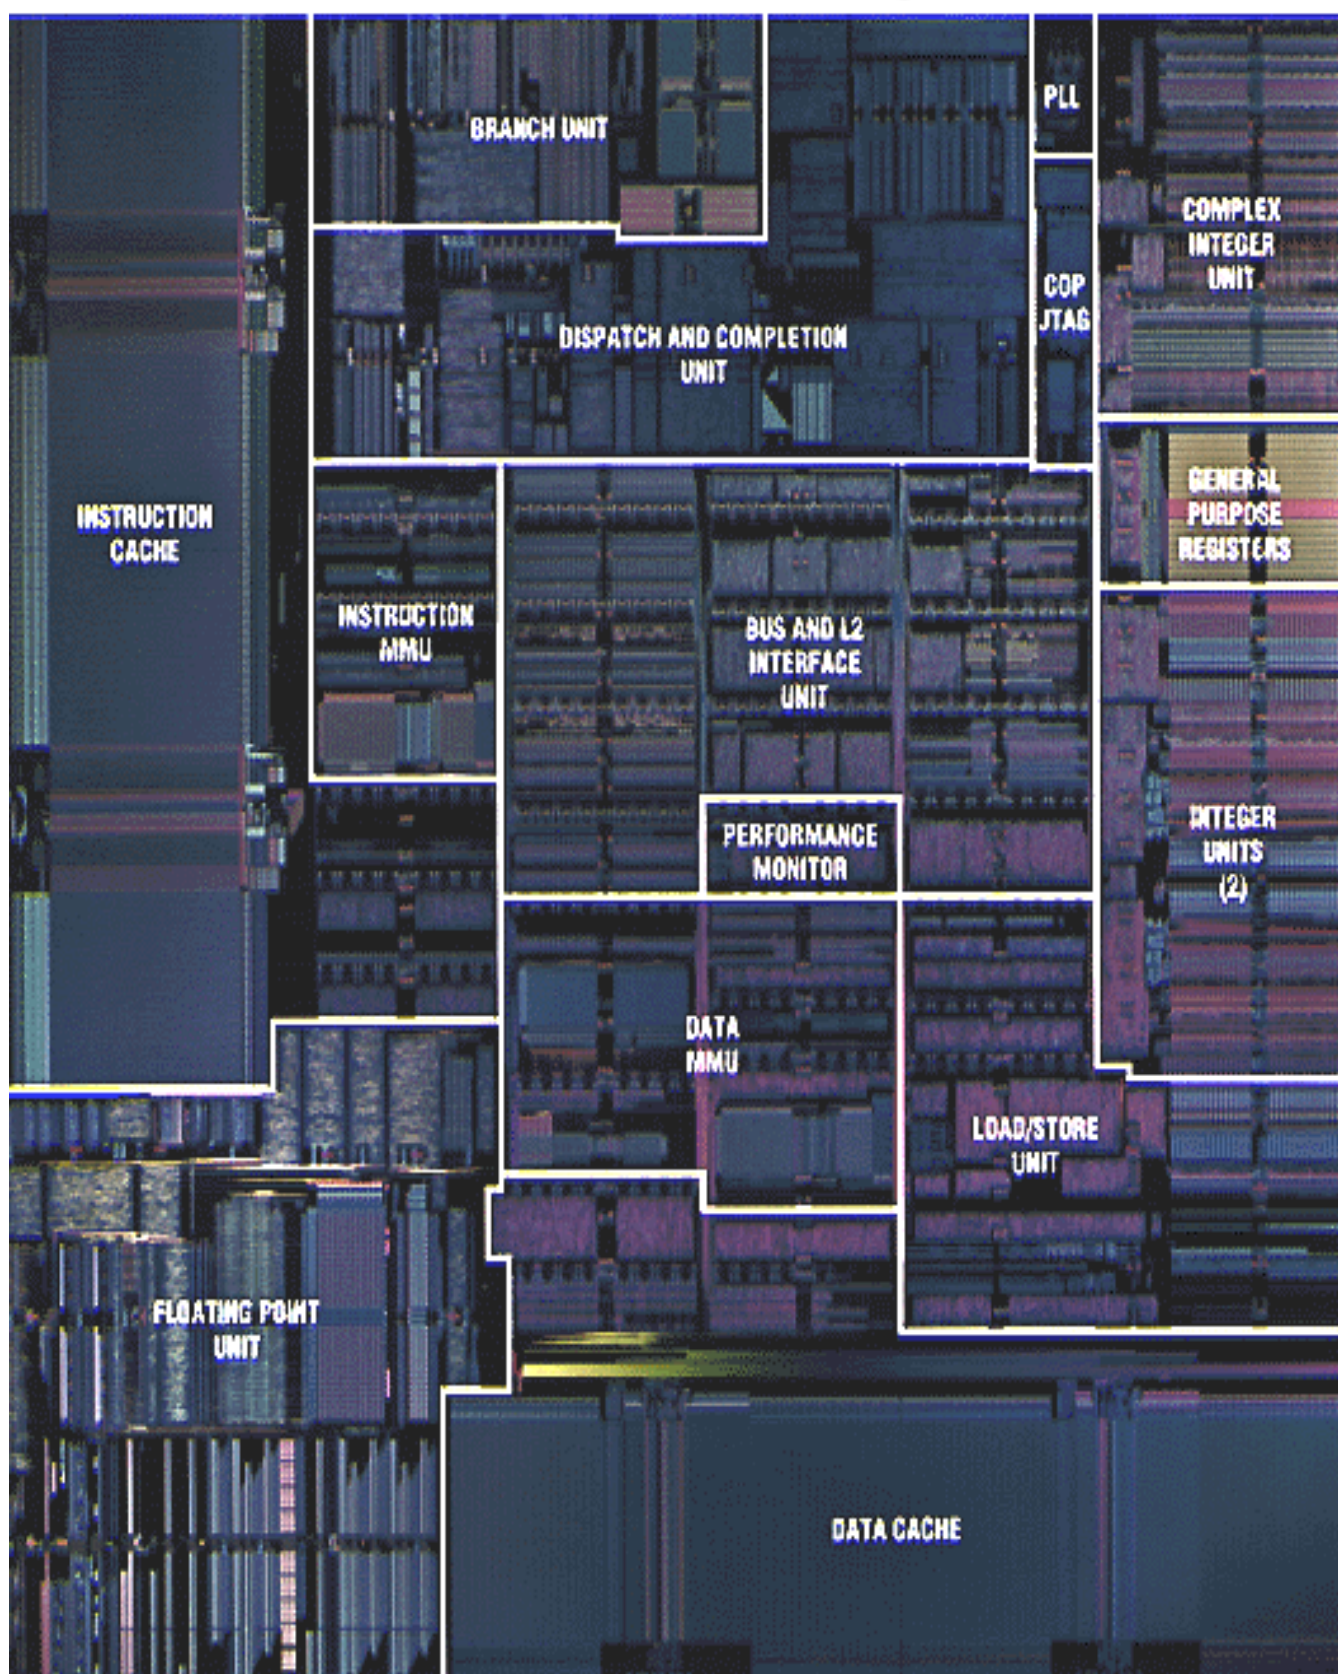


Technology

Introduction

year 1998 microprocessor PowerPC 620 7,5 Mio. transistors

Motorola's PowerPC™ 620 32/64-Bit RISC Microprocessor



Technology

Technology forecast

The 19th April 1965 issue of Electronics magazine contained an article with the title "Cramming more components onto integrated circuits" Its author, Gordon E. Moore, director, Research and Development, Fairchild Semiconductor, had been asked to predict what would happen over the next ten years in the semiconductor component industry. His article speculated that by 1975 it would be possible to cram as many as 65 000 components onto a single silicon chip about 6 mm². [Spectrum, June '97]

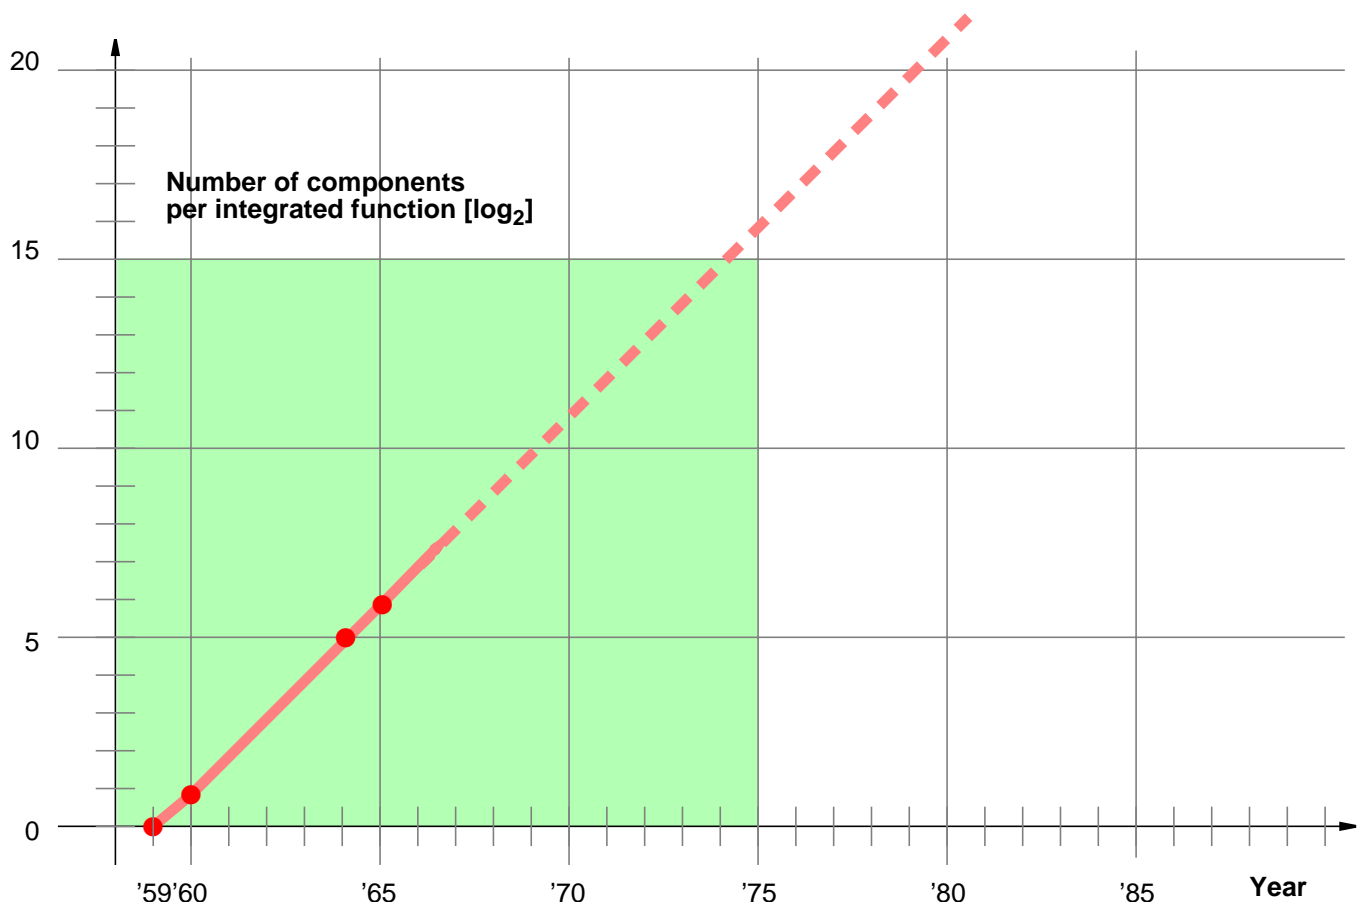
Moore based his forecast on a log-linear plot of device complexity over time.

The complexity for minimum component costs has increased at a rate of roughly a factor of two per year. Certainly over the short term this rate can be expected to continue, if not to increase. Over the longer term, the rate of increase is a bit more uncertain, although there is no reason to believe it will not remain constant for at least 10 years.

Moore's astonishing prediction was based on empirical data from just three Fairchild data points!

- first planar transistor in 1959
- few ICs in 1960 with ICs in production with 32 components in 1964
- IC with 64 components in 1965

In 1975 he revisited the slope of his function to doubling transistor count every 18 month and this stayed true until today.



Technology

Andrew (Andy) Grove

Robert Noyce

Gordon Moore



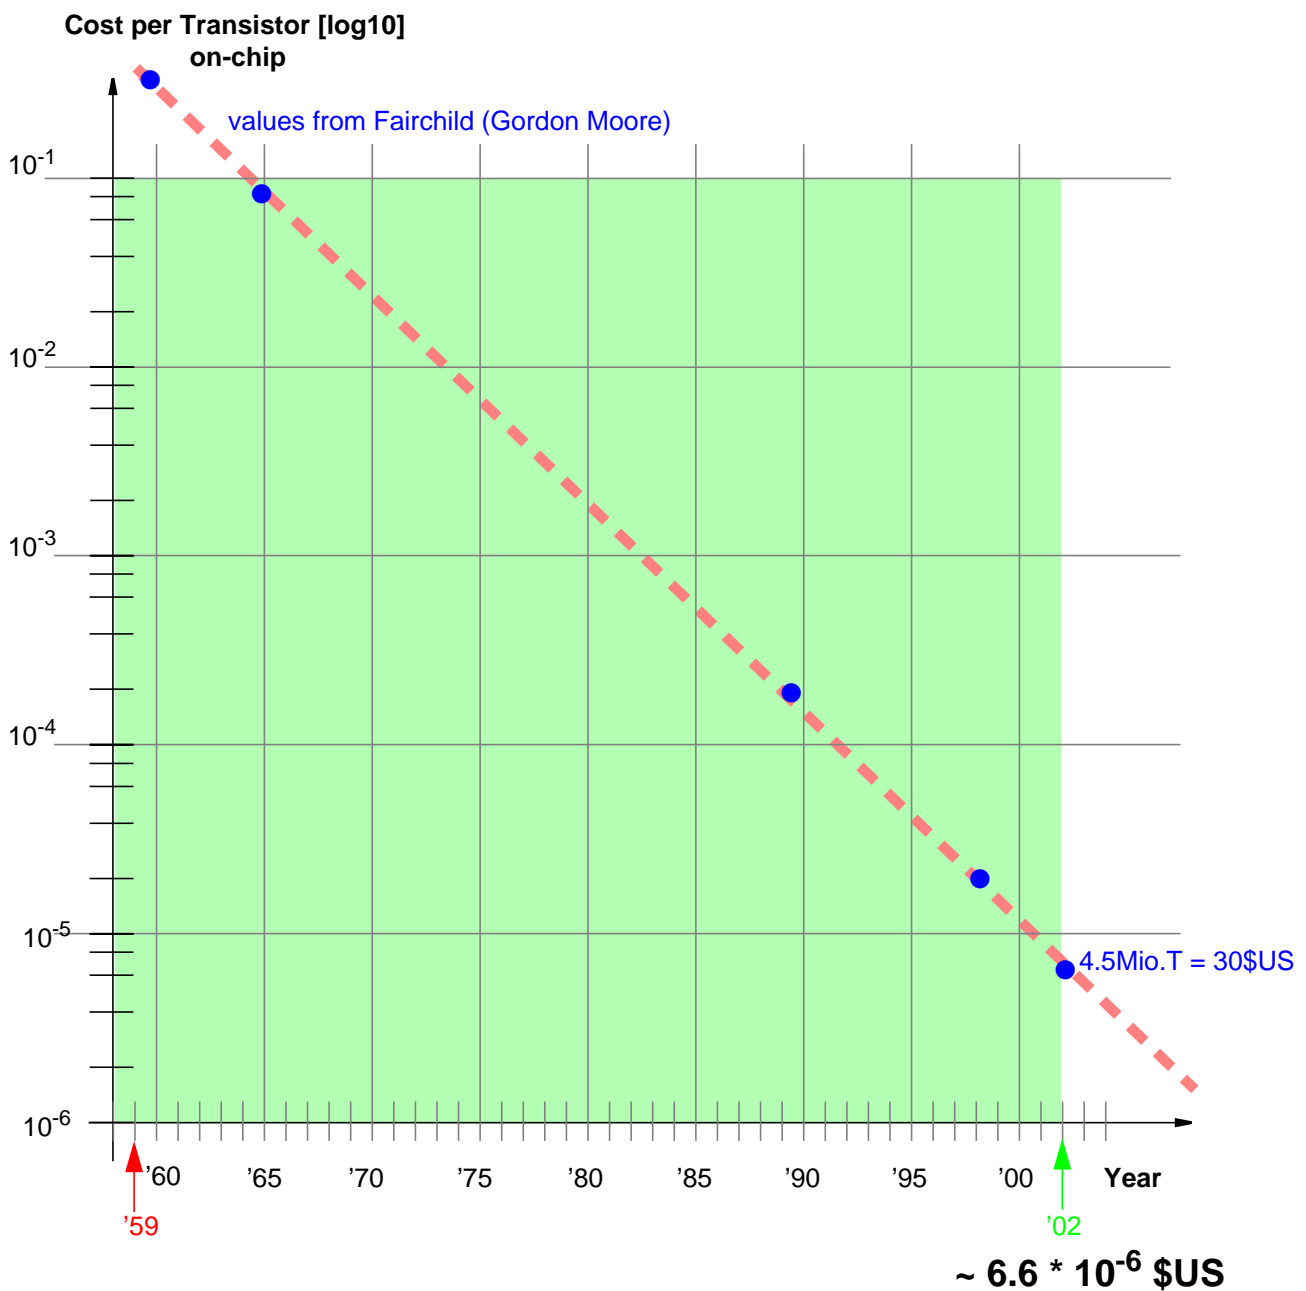
[2] Gordon Moore [right] relaxes with fellow pioneers of the electronic age: Robert Noyce [center] and Andrew Grove [left]. Moore and Noyce contributed to the development of the planar IC. Grove is now president and chief executive officer of Intel Corp.

Technology

Cost reduction due to mass fabrication

red curve for logic chips

(no memory chips)




year 2002: 4.5Mio. transistors in 0.18 μ m CMOS technology on a 5x5mm die with BGA package ~ 30 \$US = 6.6×10^{-6} \$US per transistor; standard cell design

Technology

Modern Chip Technology

Using copper for the 6 metal interconnect structure of a CMOS chip delivers lower resistance of the wires and thus increases the performance.


Gate structures can be found in the lower right part of the picture (arrow)



[1] IBM Corp.'s new CMOS 7S process for manufacturing ICs uses copper for its six levels of interconnections, and has effective transistor channel-lengths of only 0.12 μm . It is the first commercial fabrication process to use copper wires [see "The Damascus connection," p. 25].


ATOLL - ATOMIC Low Latency

- approx. 5 Mio transistors (1.4 Mio logic gates)
- Die area: 5.78mm x 5.78mm
- 0.18um CMOS (UMC, Taiwan)
- 385 staggered I/O pads
- 3 years of development



Speed Trends

Processor Memory I-O -Speed Trends



Die ständige Steigerung der Rechenleistung moderner Prozessoren führt zu einer immer größer werdenden Lücke zwischen der Verarbeitungsgeschwindigkeit des Prozessors und der Zugriffsgeschwindigkeit des Hauptspeichers. Die von Generation zu Generation der DRAMs steigende Speicherkapazität (x 4) führt auf Grund der 4-fachen Anzahl an Speicherzellen trotz der Verkleinerung der VLSI-Strukturen zu nur geringen Geschwindigkeitssteigerungen.

(not included: new DRAM technology RAMBUS)

Speed Trends

Processor Memory I-O -Speed Trends

- Results**
- a high Processor-Memory Performance Gap (increasing)
 - a very high Processor-I/O Performance Gap (increasing)

Single Processor Performance will increase

Number of processors in SMPs will increase moderately



Bus Systems will become the main bottleneck

Solution

- Transition to “switched” Systems
on all system levels
 - Memory
 - I/O
 - Network

Switched System Interconnect


Pipelining

The performance gain achieved by pipelining is accomplished by partitioning an operation F into multiple suboperations f_1 to f_k and overlapping the execution of the suboperations of multiple operations F_1 to F_n in successive stages of the pipeline [Ram77]

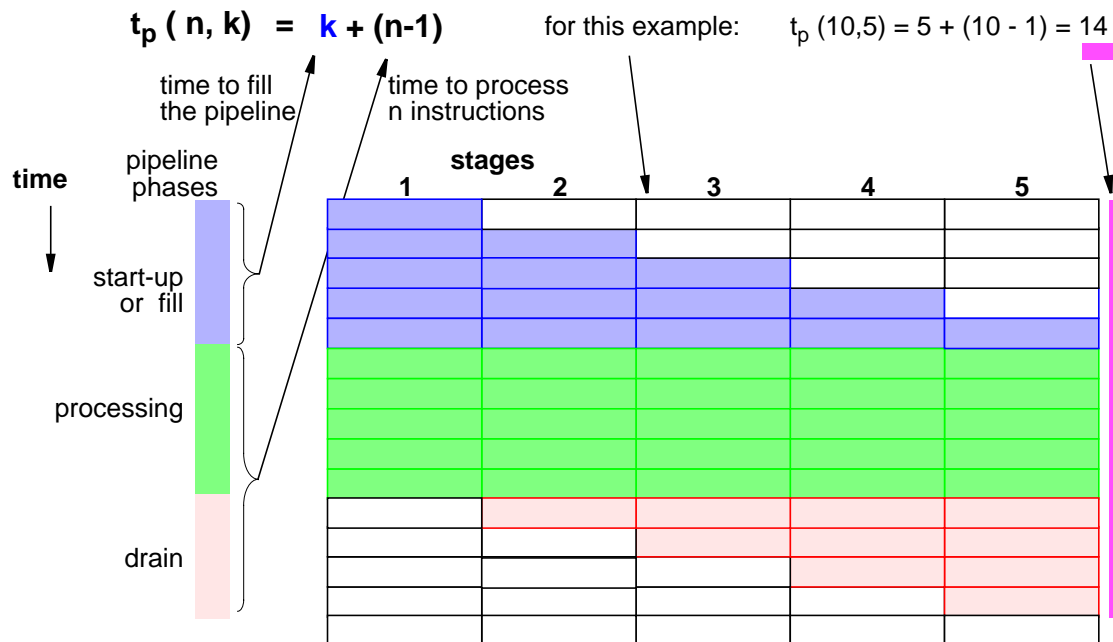
Assumptions for Pipelining

- ① the operation F can be partitioned
- ② all suboperations f_i require approximately the same amount of time
- ③ there are several operations F the execution of which can be overlapped
- ④ the execution time of the suboperations is long compared with the register delay time

Linear Pipeline with k Stages



Pipelined Operation Time



Durchsatz Throughput

$$TP(n, k) = \frac{\text{number of operations}}{t_p(n, k)} = \frac{n}{k + (n-1)} \left[\frac{\text{operations}}{\text{time unit}} \right]$$

Gewinn Gain

$$S(n, k) = \frac{\text{scalar execution time}}{\text{pipelined execution time}} = \frac{n * k}{k + (n-1)} \quad \lim_{n \rightarrow \infty} S \rightarrow k$$

Effizienz Efficiency

initiation rate, latency

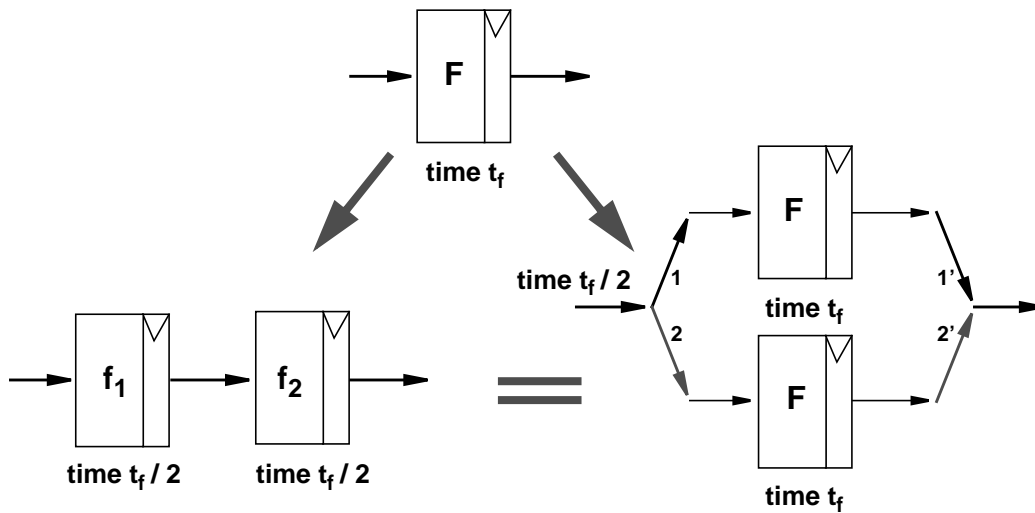
$$E(n, k) = \frac{1}{k} * S(n, k) = \frac{n * k}{k * (k + (n-1))} = \frac{n}{k + (n-1)}$$

Pipeline Interrupts

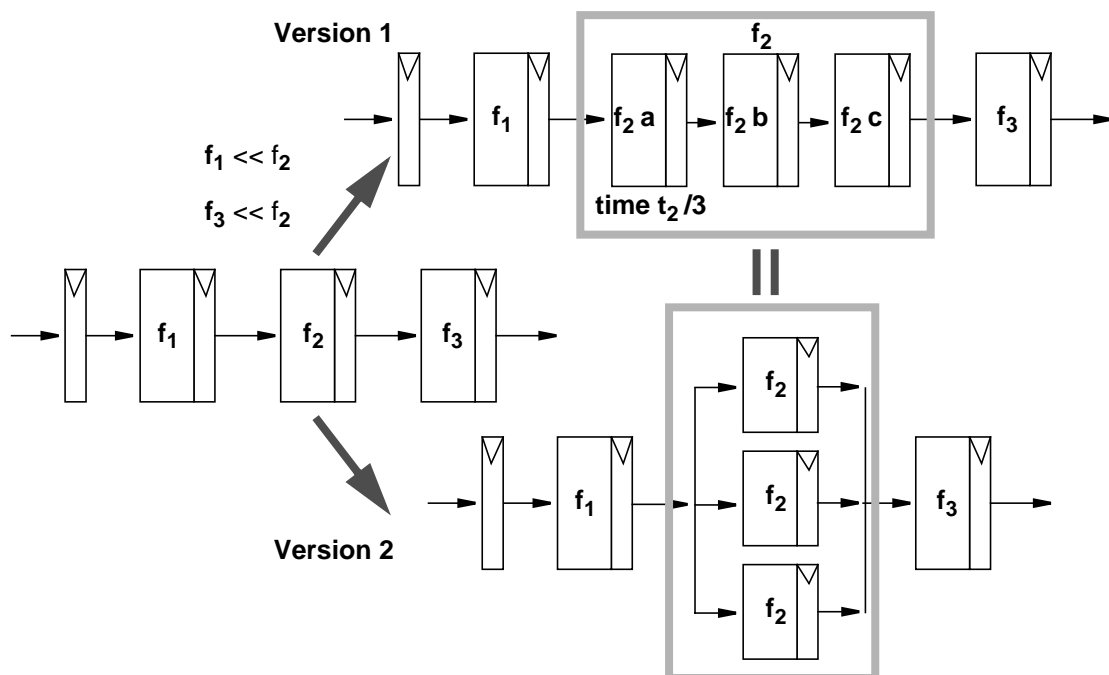
- data dependencies
- control-flow dependencies
- resource dependencies

Assumptions for Pipelining

- ① the operation F can be partitioned



- ② all suboperations f_i require approximately the same amount of time



Assumptions for Pipelining

- ③ there are several operations F
the execution of which can be overlapped

If there is a discontinuous stream of operations F owing to a conflict, bubbles are inserted into the pipeline. This reduces the performance gain significantly.


A typical example of this is the control dependency of the instruction pipeline of a processor. Here, each conditional branch instruction may disrupt the instruction stream and cause $(k-1)$ bubbles (no-operations) in the pipeline, if the control flow is directed to the nonpredicted path.

- ④ the execution time of the suboperations is long
compared with the register delay time

Assuming a division of the operation F into three suboperations f_1, f_2, f_3 , and also no pipelining, the operation F can be executed in the time:

$$t(F) = t_{f1} + t_{f2} + t_{f3}$$

Introduction of registers



$$t(F) = \underbrace{(\max(t_{fi}) + t_{co} + t_{su})}_{t_{cyc}} * \underbrace{3}_{k \text{ stages}} = 3 * \max(t_{fi}) + \underbrace{3 * (t_{co} + t_{su})}_{\text{register delay time}}$$

$$t_{cyc} = \max(t_{fi}) + t_{co} + t_{su} \quad f_{cyc} = 1 / t_{cyc}$$

The registers are introduced behind each function of the suboperation and this creates the pipeline stages. Placing the register at the output (not at the input!!!) makes suboperation stages compatible with the definition of state machines, which are used to control the pipeline.

Data Flow Dependencies

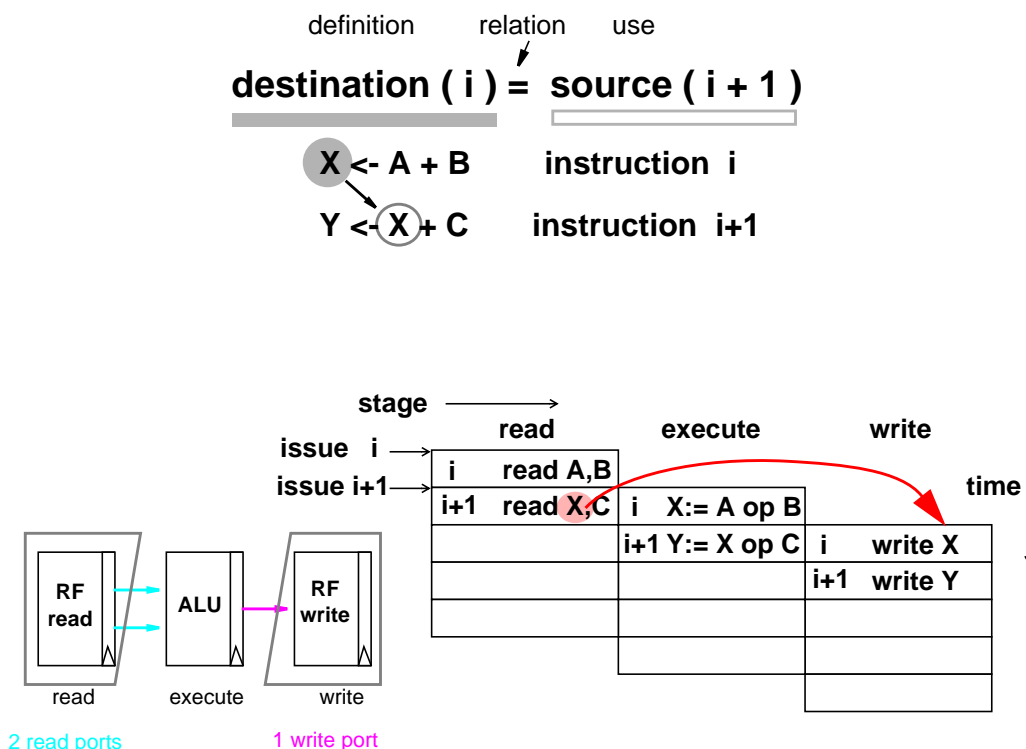
Three different kinds of data flow dependency hazards can occur between two instructions. These data dependencies can be distinguished by the definition-use relation.

- flow dependency read after write RAW
- anti-dependency write after read WAR
- output dependency write after write WAW

Basically, the only real dependencies in a program are the data flow dependencies, which define the logical sequence of operations transforming a set of input values to a set of output values (see also Data Flow Architectures). Changing the order of instructions must consider these data flow dependencies to keep the semantic of a program.

The **Flow Dependency** read after write (RAW)

This data dependency hazard can occur e.g. in a pipelined processor where values for a second instruction $i+1$ are needed in the earlier stages and have not yet been computed by the later stages of instruction i . The hazardous condition arises if the dependent instructions are closer together than the number of stages among which the conflicting condition arises.




To avoid the hazard, the two instructions must be separated in time within the pipeline. This can be accomplished by **inserting bubbles** (NOPs) into the pipeline (simply spoken: by waiting for i to complete) or building a special hardware (**hardware interlock**), which inserts NOPs automatically.

Data Flow Dependencies

The **Anti-Dependency** write after read (WAR)

This dependency can only arise in the case of out-of-order instruction issue or out-of-order completion. The write back phase of instruction $i+1$ may be earlier than the read phase of instruction i . Typical high-performance processors do not use out-of-order issue. This case, then, is of less importance. If the compiler reorders instructions, this reordering must preserve the semantics of the program and take care of such data dependencies.

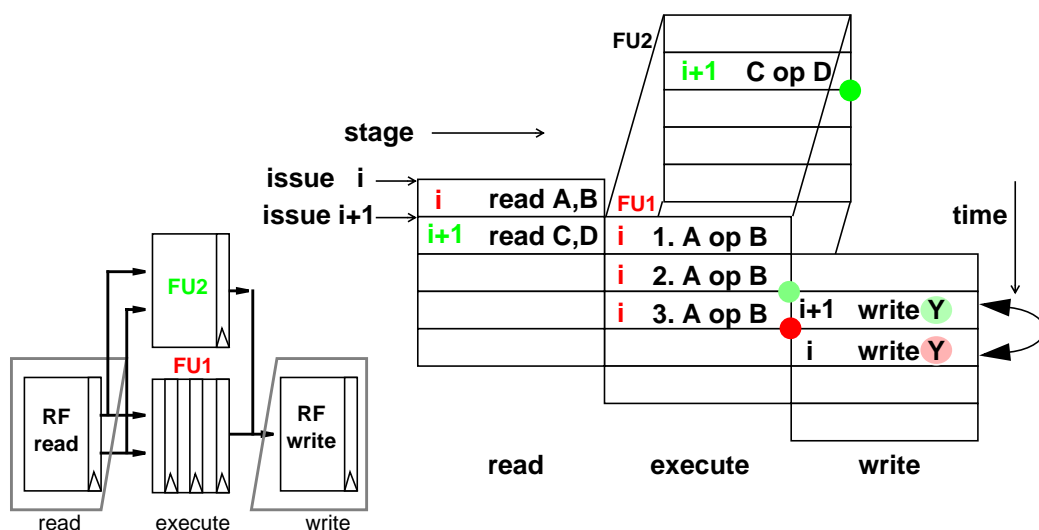

$$\text{source}(i) = \text{destination}(i+1)$$



The **Output Dependency** write after write (WAW)

The result y of the first instruction i will be written back to the register file later than y of instruction $i+1$ because of a longer pipeline for the division. This reverses the sequence of writes i and $i+1$ and is called out-of-order completion.

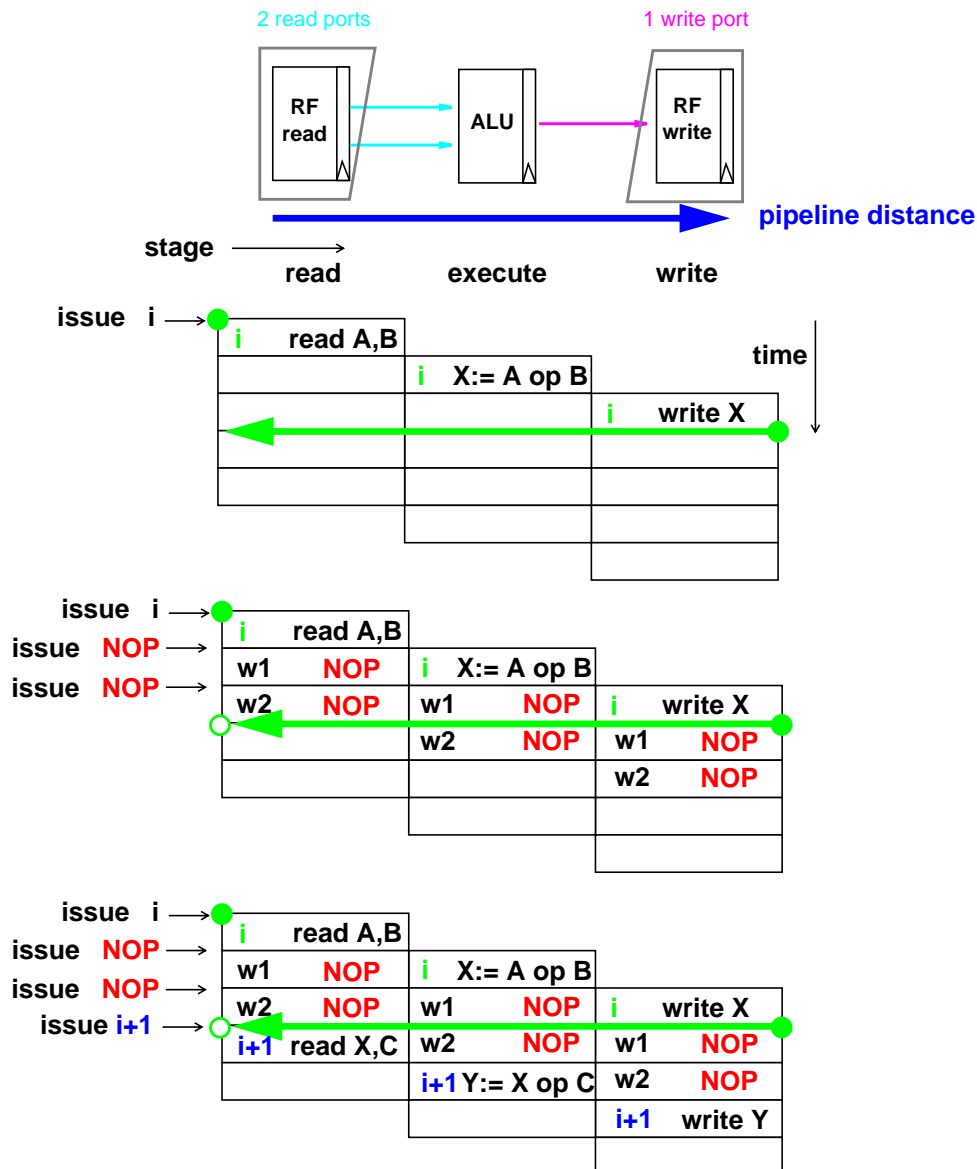
$$\text{destination}(i) = \text{destination}(i+1)$$



Data Flow Dependencies

Inserting Bubbles

We must avoid the RAW-hazard in the pipeline, because reading X before it is written back from the previous instruction reads an old value of X and thus destroys the semantic of the two sequential instructions. The sequential execution model (von Neumann architecture) assumes that the previous instruction is completed before the execution advances to the next instruction.



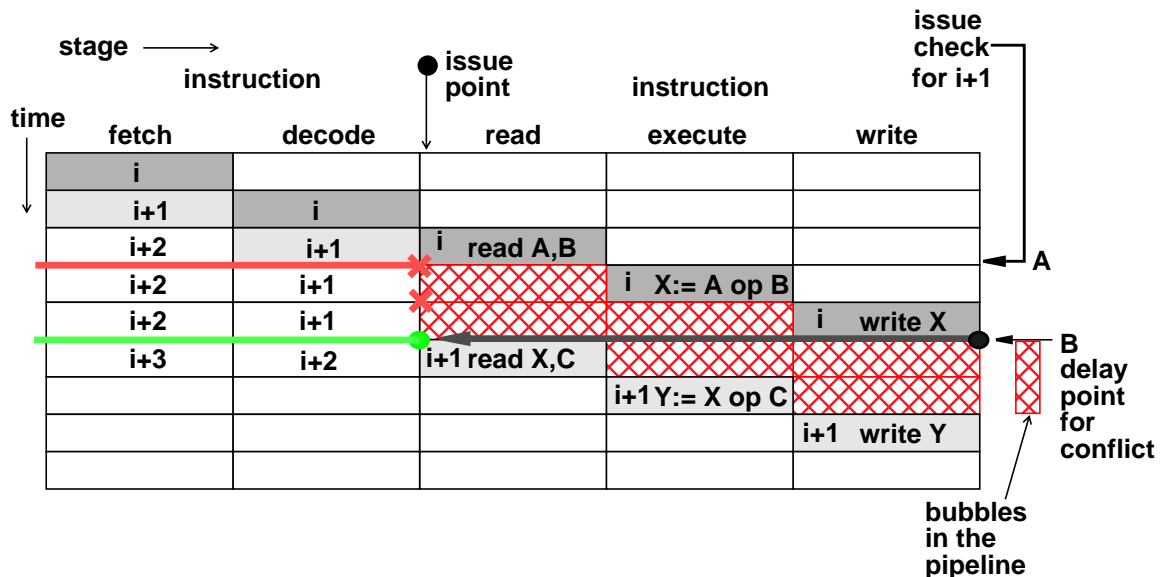
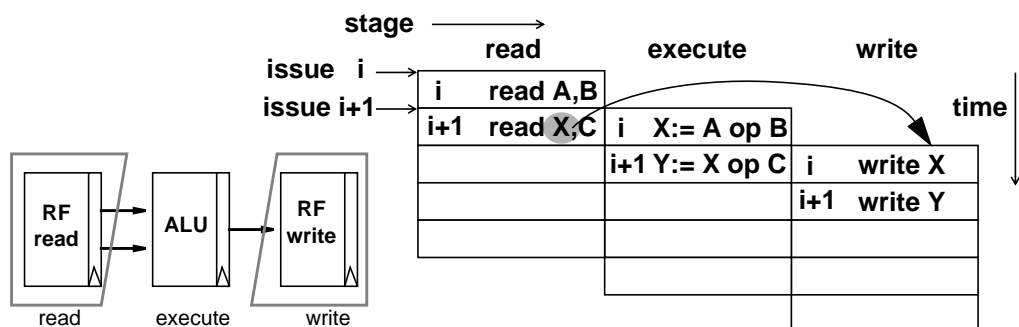
The conflicting stages are RF read and the RF write, which have a distance of 3 pipeline stages. Reading the correct value for X from the register file requires the insertion of 2 NOPs in between the two conflicting instructions. This delays the instruction *i+1* by 2 clocks which then removes the RAW-hazard. The compiler (or programmer) is responsible for detecting the hazard and inserting NOP into the instruction stream.

The 2 bubbles in the pipeline can be utilized by other useful instructions independent from *i* and *i+1* (see instruction scheduling).

Data Flow Dependencies

Hardware Interlock

A hardware mechanism for detecting and avoiding the RAW-Hazard is the interlock hardware. The RAW-Hazard condition is detected by comparing the source fields of instruction $i+1$ (and $i+2$) with the destination field of instruction i , until i is written back to the register file and thus the instruction is completely executed.

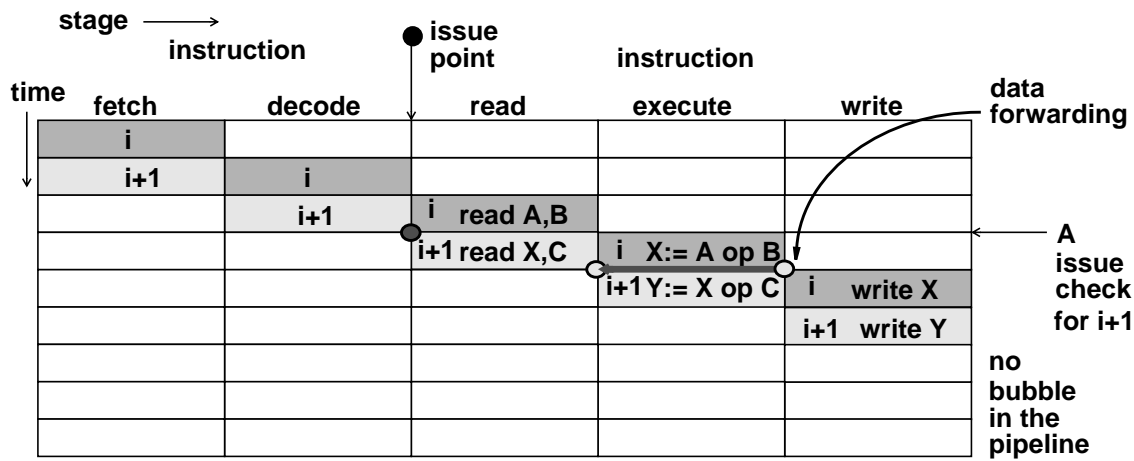
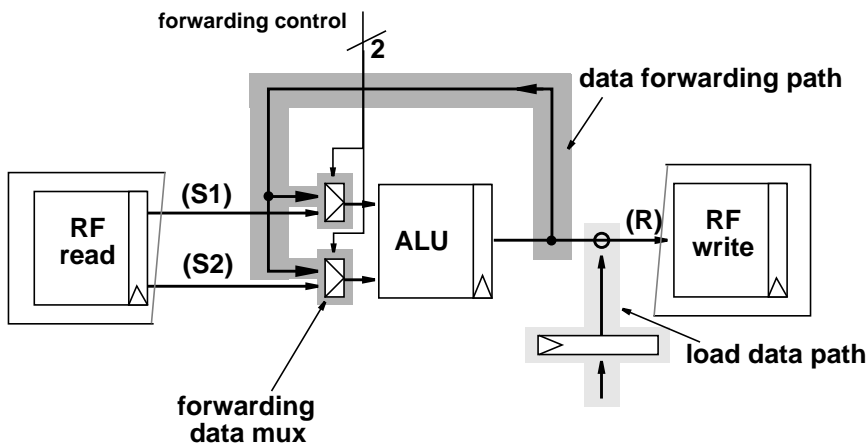


The hardware interlock detects the RAW-Hazard and delays the issue of instruction $i+1$ as long as the write back of instruction i is completed.

The hardware mechanism doesn't need additional NOPs in the instruction stream. The bubbles are inserted by the hardware.

Nevertheless the produced bubbles can be avoided by scheduling useful and independent instructions in between i and $i+1$ (see also instruction scheduling).

Data Forwarding



Control Flow Dependencies

The use of pipelining causes difficulties with respect to control-flow dependencies. Every change of control flow potentially interrupts the smooth execution of a pipeline and may produce bubbles in it. One of the following instructions may be responsible for them:

- conditional branch instruction
- jump instruction
- jump to subroutine
- return from subroutine

The bubbles (no operations) in the pipeline reduce the gain of the pipeline thus reducing performance significantly. There are two causes of bubbles. The first one is a data dependency in the pipeline itself, the branch condition being evaluated a pipeline stage later than needed by the instruction fetch stage to determine the correct instruction path. The second one is the latency for accessing the instruction and the new destination for the instruction stream.

Pipeline Utilization

$$U = \frac{n}{n+m} = \frac{1}{1 + \frac{m}{n}}$$

n number of useful instructions
m number of no operations (bubbles)

$$m = \underbrace{b \cdot (1-p) \cdot N_b}_{\text{no ops caused by branches}} + \underbrace{b \cdot q \cdot N_o}_{\text{no ops caused by the latency of branch target fetches}}$$

no ops caused by branches

no ops caused by the latency of branch target fetches

b number of branches
p probability of correct guesses
N_b penalty for wrong guess
q frequency of other causes
N_o penalty for other causes

$$U = \frac{1}{1 + \frac{b}{n} \cdot (1-p) \cdot N_b + \frac{b}{n} \cdot q \cdot N_o}$$

Control Flow Dependencies

Reduction of Branch Penalty

The 'no' operations (NOPs) in the pipeline can be reduced by the following techniques:

reduction of branch penalty N_b

- forwarding of condition code
- fast compare
- use of delay slots
- delayed branch
- delayed branch with squashing

increase of p

- static branch prediction
- dynamic branch prediction
- profiled branch prediction

reduction of instruction fetch penalty N_o

- instruction cache
- improvement of instruction prefetch
- branch target buffer
- alternate path following

avoid branches

- predication of instructions

Control Flow Dependencies

Effect of Branches

Before we start to analyse the effects of branching we should introduce the basic terms and have a closer look to the internals of the instruction fetch unit.

Definition : *Branch successor: The next sequential instruction after a (conditional) branch is named the branch successor.*

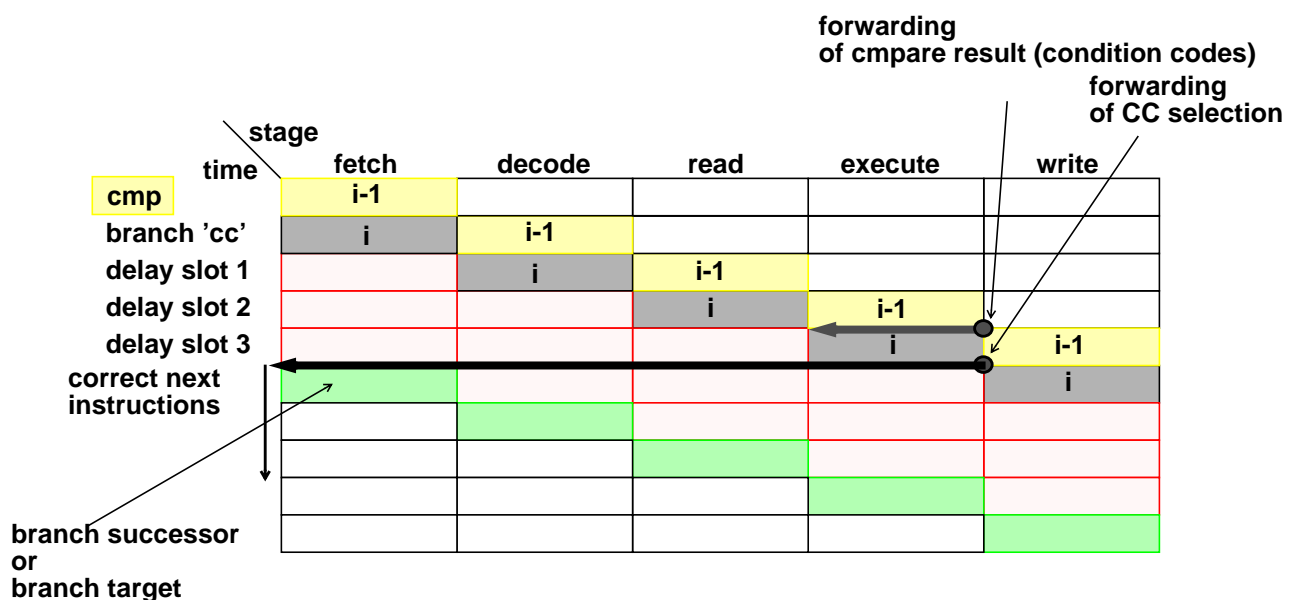
In this case, it is assumed the branch is not taken. The operation of fetching the next instruction in the instruction stream requires to add the distance from the actual instruction to the next instruction to the value of the actual program counter (PC).

Definition : *Branch target: The instruction to be executed after a branch is taken is called a branch target.*

The operation required to take the branch is to add the branch offset from the instruction word to the PC and then fetch the the branch target instruction from the instruction memory.

The number of pipeline cycles wasted between a branch taken and the resumption of instruction execution of the branch target is called a delay slot. Depending on the pipeline the number of delay slots b may vary between 0 and $k-1$.

In the following stage-time diagram we assume that the instruction fetch can immediately continue when the branch direction is defined by the selection of the condition code, which is forwarded to the instruction fetch stage.


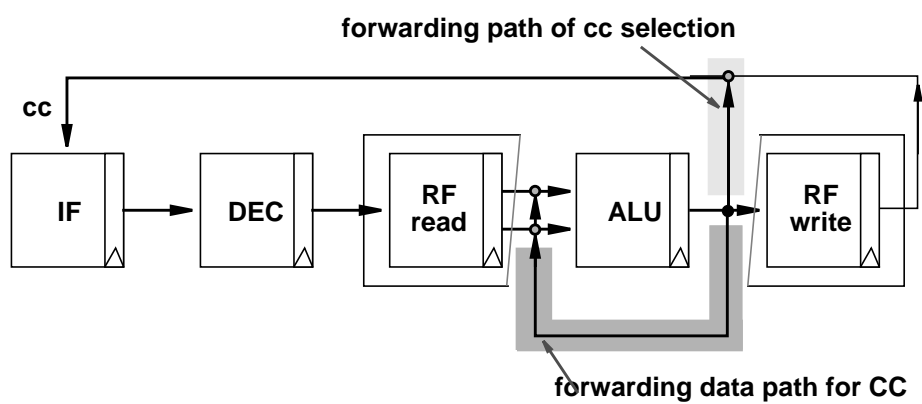


Control Flow Dependencies

Instruction fetch stage

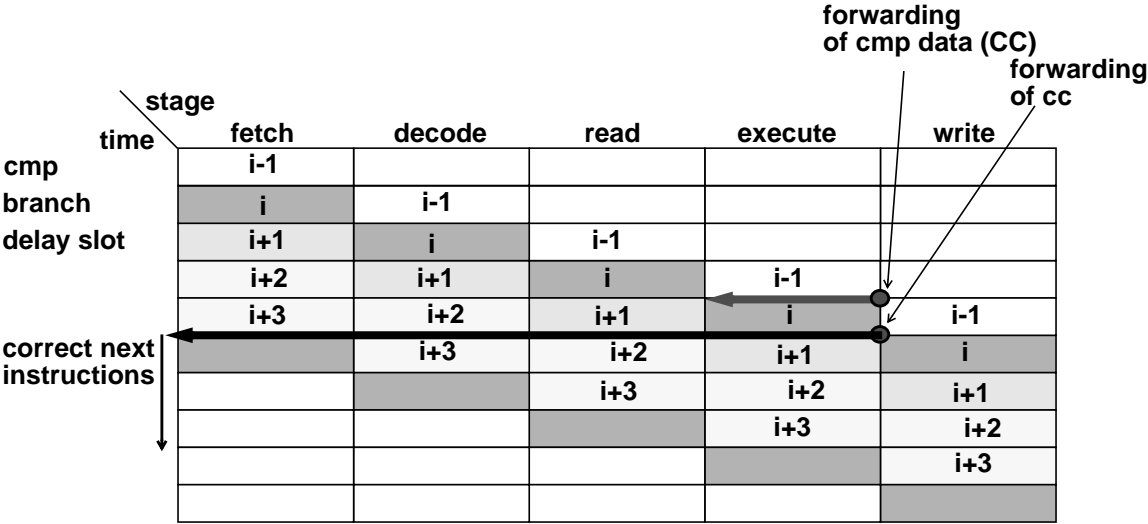
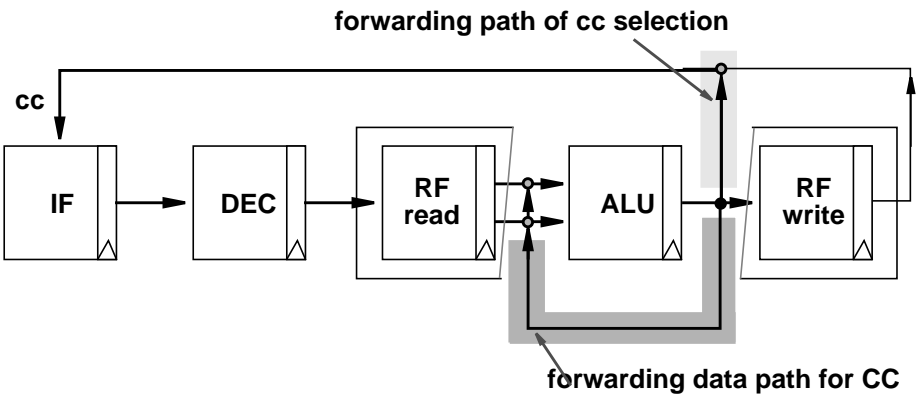
cmp R1,R2 -> CC calculate all reasonable bits

bra CC, offset change the control flow
depending on the selection of CC bits
calculate next PC



numbering is used to identify the instructions in the flow of processor code

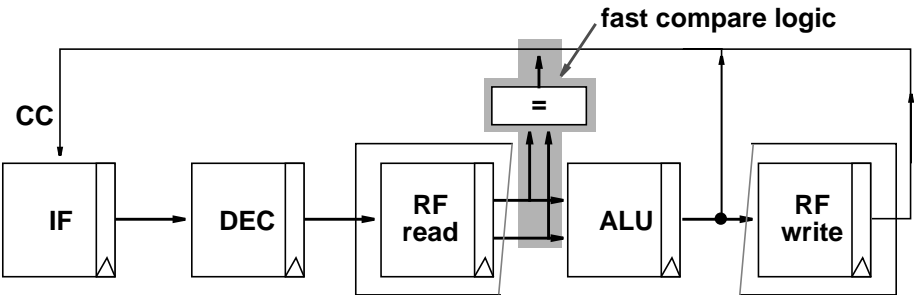
Forwarding of the Condition Code



Control Flow Dependencies

Fast Compare

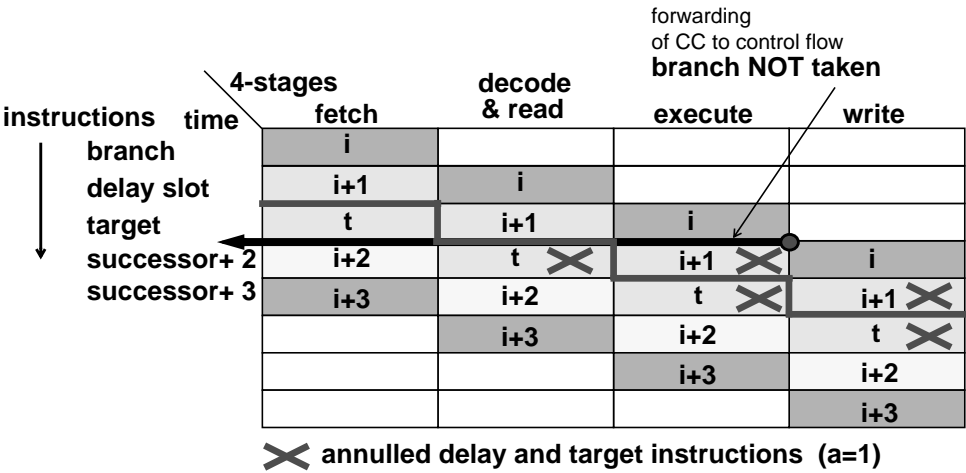
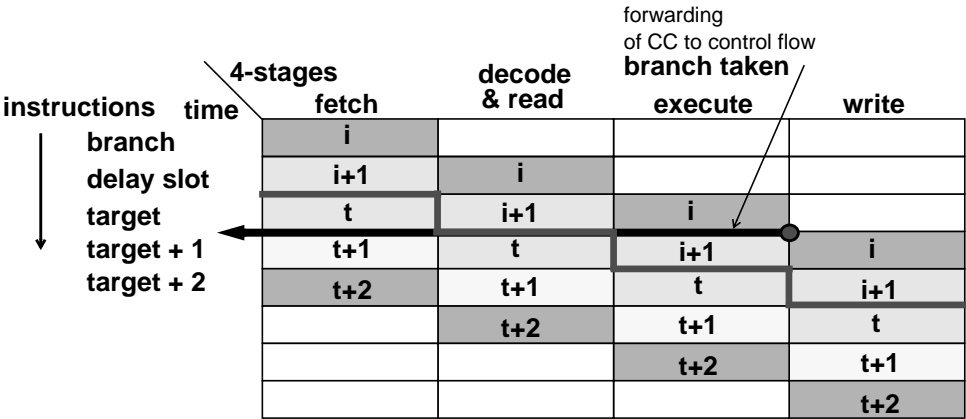

simple tests for equal, unequal, <0, <=0, >0, >=0, =0



forwarding of fast cmp data

| stage | fetch | decode | read | execute | write |
|---------------------------|-------|--------|------|---------|-------|
| cmp | i-1 | | | | |
| branch | i | i-1 | | | |
| delay slot | i+1 | i | i-1 | | |
| target | i+2 | i+1 | i | i-1 | |
| correct next instructions | | i+2 | i+1 | i | i-1 |
| | | | i+2 | i+1 | i |
| | | | | i+2 | i+1 |
| | | | | | i+2 |
| | | | | | |
| | | | | | |

Delayed Branch with Squashing



If the anullment bit is asserted, the delay slot instruction will be annulled, when the branch decision was predicted falsely.

This enables the compiler to schedule instructions from before and after the branch instruction into the delay slot.

Branch Prediction

Two different approaches can be distinguished here:


- static-branch prediction
- dynamic-branch prediction

The static-branch information associated with each branch does not change as the program executes. The **static-branch prediction** utilizes compile time knowledge about the behaviour of a branch. The compiler can try to optimize the ordering of instructions for the correct path.

Four prediction schemes can be distinguished:

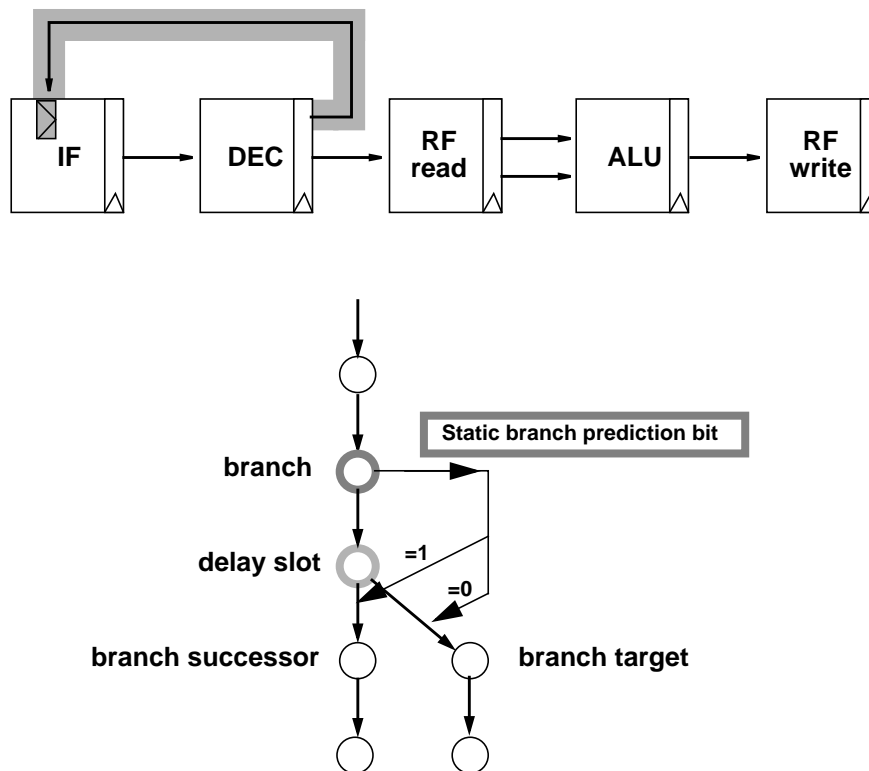
- branch not taken
- branch taken
- backward branch taken
- branch direction bit

built-in static prediction strategies



Static Branch Prediction

The hardware resources required for static branch prediction are a forwarding path for the prediction bit, and a logic for changing the instruction fetch strategy in the IF stage. This prediction bit supplied by the output of the instruction decode stage decides which path is followed. The prediction bit allows execution of instructions from the branch target or from the successor path directly after the delay slot instruction.



Profiled Branch Prediction

The **profiled branch prediction** is a special case of the static prediction. The guess of the branch direction is based on a test run of the program. This profile information on branch behaviour is used to insert a static branch prediction bit into the branch instruction that has a higher proportion of correct guesses.


Dynamic Branch Prediction

Dynamic branch prediction uses information on branch behaviour as the program executes. No initial dynamic information exists before the program starts executing.

The dynamic information is normally associated with a specific branch. The association can be realized by using the instruction address of the branch as the unique identifier of the branch.

The dynamic prediction depends on the past behaviour of this branch and is stored in a table addressed by the branch instruction address. A unique addressing would need the whole address as an index, usually 32 bit in length. This length prohibits direct use of this address as an index to the state table.


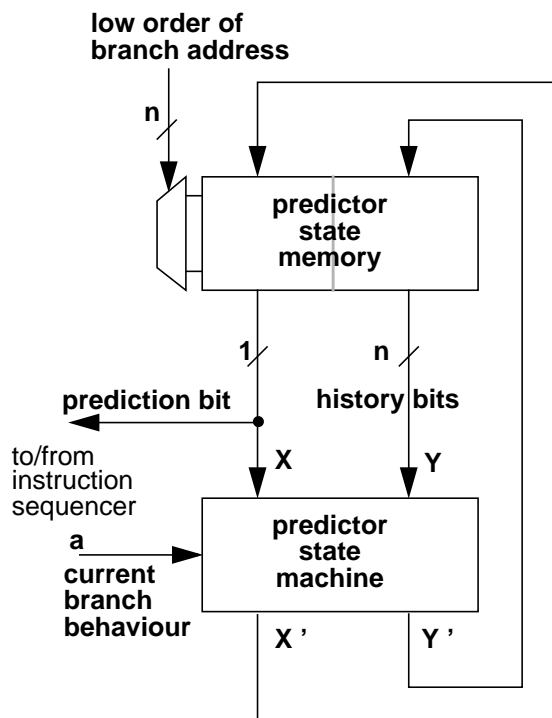
simple branch predictor using only one history bit



a current branch behaviour

X prediction bit

Y history bits

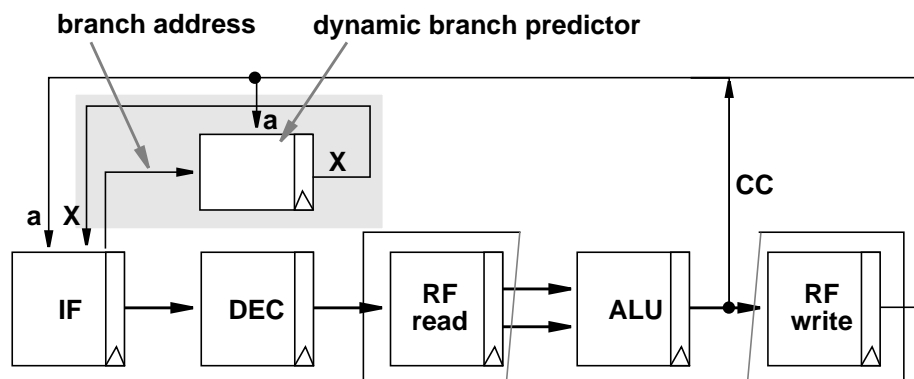


a current branch behaviour
taken a
not taken /a

X prediction bit
take 1
take not 0

Y history bit
last branch taken 1
last branch not taken 0

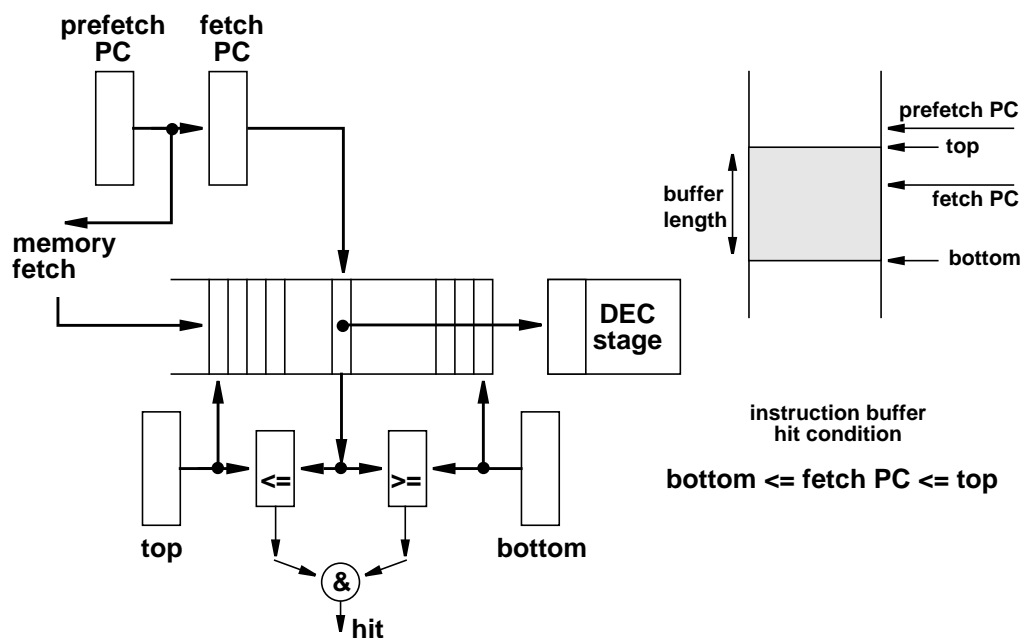
Dynamic Branch Prediction



Reduction of Instruction Fetch Penalty

instruction cache

instruction buffer



Hardware Resources for an Instruction Buffer

Reduction of Instruction Fetch Penalty

Branch Target Buffer

